

# Programming in Java

## Threads and multi-threading — the Fork/Join API

KLM

Department of Computer Science and Information Systems  
Birkbeck, University of London

`keith@dcs.bbk.ac.uk`



# High Level Concurrency Objects

As concurrency was seen as too low level in previous versions of Java some new classes were introduced in

**Lock objects** support locking idioms that simplify many concurrent applications

**Executors** define a high-level API for launching and managing threads

**Concurrent collections** make it easier to manage large collections of data, and can greatly reduce the need for synchronisation = good!

**Atomic variables** have features that minimise synchronisation and help avoid memory consistency errors

**ThreadLocalRandom** provides efficient generation of pseudo-random numbers from multiple threads

# Executors — ForkJoin I

- The use of `ForkJoinPool` and `RecursiveTask`
- Using the fork/join framework is straight forward
- The first step is to write some code that performs a segment of the work
- Your code should look similar to this:

```
if (my portion of the work is small enough)
    do the work directly
else
    split my work into two pieces
    invoke the two pieces and wait for the results
```

# Executors — ForkJoin II

- wrap this code as a `ForkJoinTask` subclass
- typically as one of its more specialised types
  - `RecursiveTask` (which can return a result) or
  - `RecursiveAction`
- After your `ForkJoinTask` is ready
- create one that represents all the work to be done and pass it to the `invoke()` method of a `ForkJoinPool` instance

# Executors — ForkJoin III

```
package helloWorld;
import java.util.concurrent.ForkJoinPool;
import java.util.concurrent.RecursiveTask;

public class Incrementor extends RecursiveTask<Integer>{
    ...
}
```

- Next you need to create a class that extends `RecursiveTask`
- A `RecursiveTask` is like a `Thread`
- But you can retrieve a value from it after it finishes
- Supply a type parameter (`Integer`, in this example) to specify the kind of value you want

# Executors — ForkJoin IV

```
package helloWorld;
import java.util.concurrent.ForkJoinPool;
import java.util.concurrent.RecursiveTask;

public class Incrementor extends RecursiveTask<Integer>{
    public static ForkJoinPool fjPool = new ForkJoinPool();
    ...
}
```

- Then, create a `ForkJoinPool`
- Create only one of these, and it should be `static`
- A `ForkJoinPool` is a pool of `Thread` objects.
- Java will take care of all the thread allocation and deallocation for you

# What did that code do? I

- Is that any different from a standard method call?
- Answer: yes, `compute`, which replaces `run`, runs as a `RecursiveTask`
- So what?

# What did that code do? II

- `RecursiveTasks` are, so-called, `lightweight` threads (you can have 32K of them) which can happily `fork` and `join`
- `public final ForkJoinTask<V>fork()` arranges to asynchronously execute this task  
Note: any given task should only be forked once
- `public final V join()`, unlike the `Thread join`, this returns a result



# Another example I

Look at the code...

# Another example II

- The call to `fjPool.invoke` creates an instance of this `Sum` class, which is a `RecursiveTask`, which gets things started
- `invoke` should not be called from within a `RecursiveTask` or `RecursiveAction`
- It should only be called from sequential code

# Atomic variables

- The `java.util.concurrent.atomic` package defines classes that support atomic operations on single variables
- All classes have `get` and `set` methods that work like reads and writes on volatile variables
- a `set` has a `happens-before` relationship with any subsequent `get` on the same variable
- The atomic `compareAndSet` method also has these memory consistency features, as do the simple atomic arithmetic methods that apply to integer atomic variables

Example...

We've barely scratched the surface...

So please do follow up the links on the Moodle site

# Questions

